

docker



Agenda

- Containers VS Virtual Machines
- Container technology
- Docker architecture
- Docker Images
- Docker Containers
- Install Docker
- Images Commands
- Containers Commands
- Conclusion



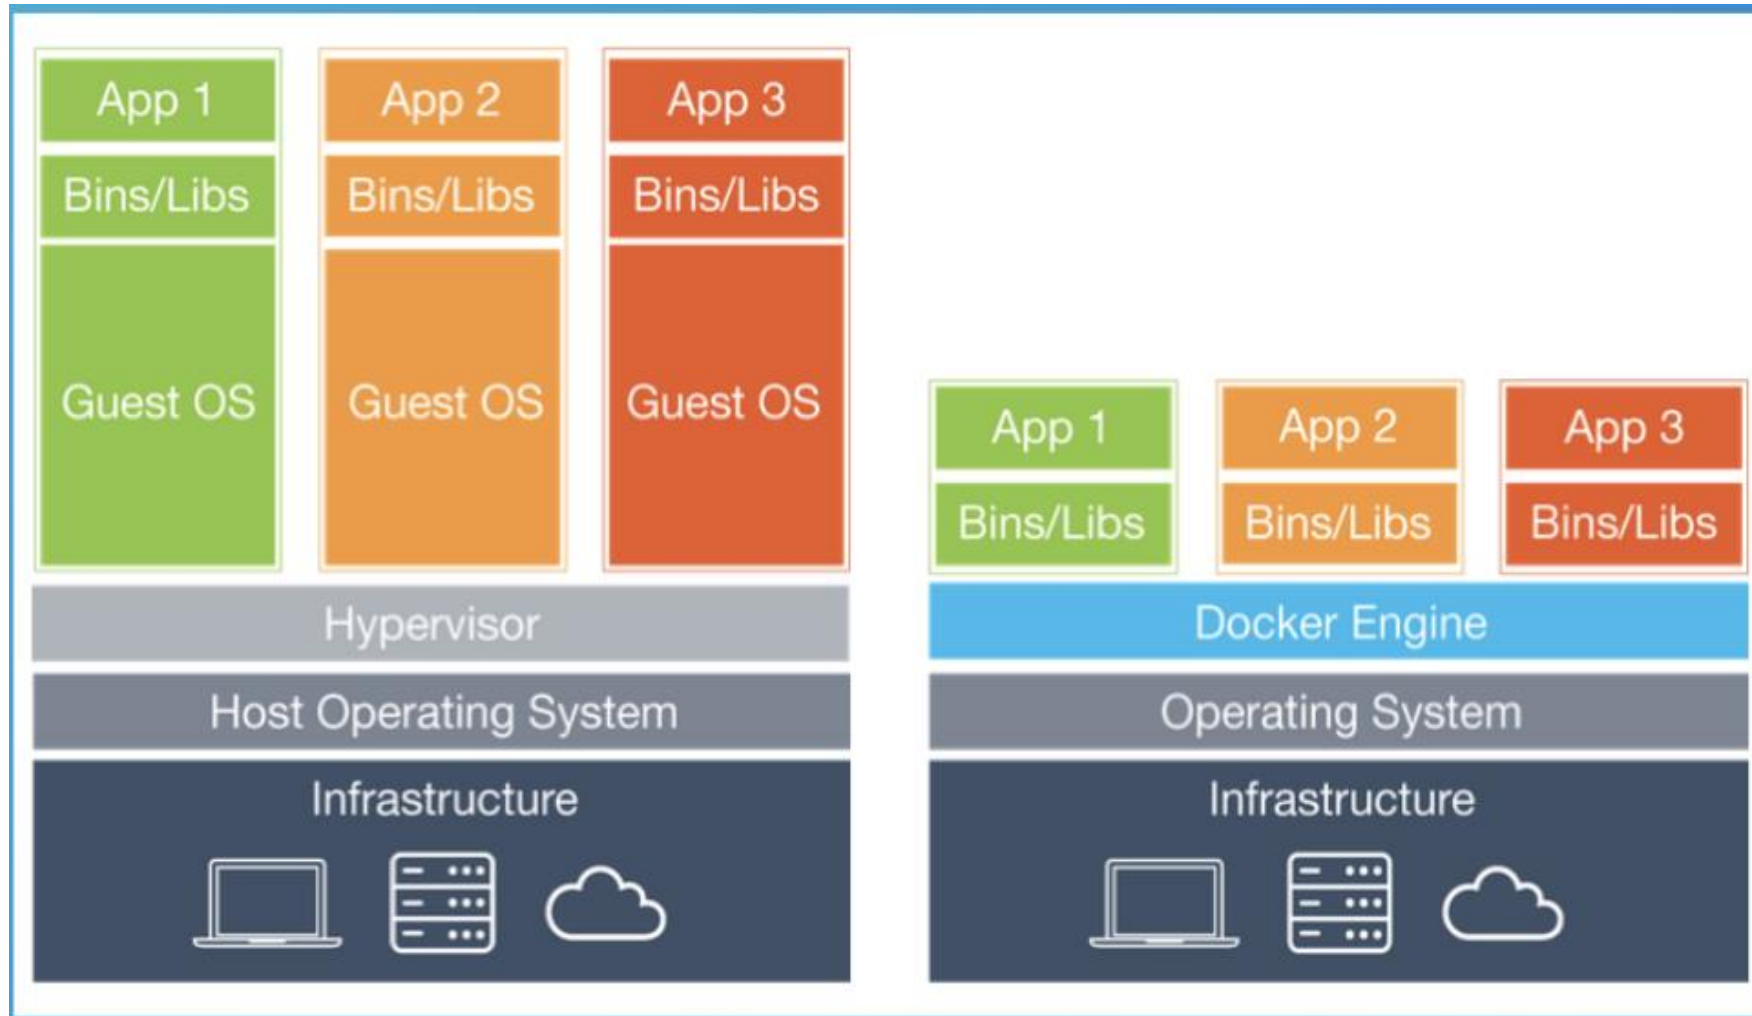
Agenda

- Docker file basics & Sample
- Building Docker file, tagging and creating image
- Logging to your Docker repo from CLI
- Pushing/Pulling image to docker repo

Day 2

- Docker Volume
- Docker Networking

Virtual Machines VS Containers





Virtual Machines VS Containers

VMs	Containers
Heavyweight	Lightweight
Limited performance	Native performance
Each VM runs in its own OS (More secure and isolated)	All containers share the host OS
Slower Startup	Faster Startup
More memory is required	Fewer memory is required

Container technology



Namespaces

A namespace **isolates** specific system resources usually visible to all processes. Inside a namespace, only processes that are members of that namespace can see those resources. Namespaces can include resources like network interfaces, the process ID list, mount points, IPC resources, and the system's host name information.

Control groups (cgroups)

Control groups partition sets of processes and their children into groups to manage and **limit the resources they consume**. Control groups place restrictions on the amount of system resources processes might use. Those restrictions keep one process from using too many resources on the host.



Container History cont'd

Seccomp

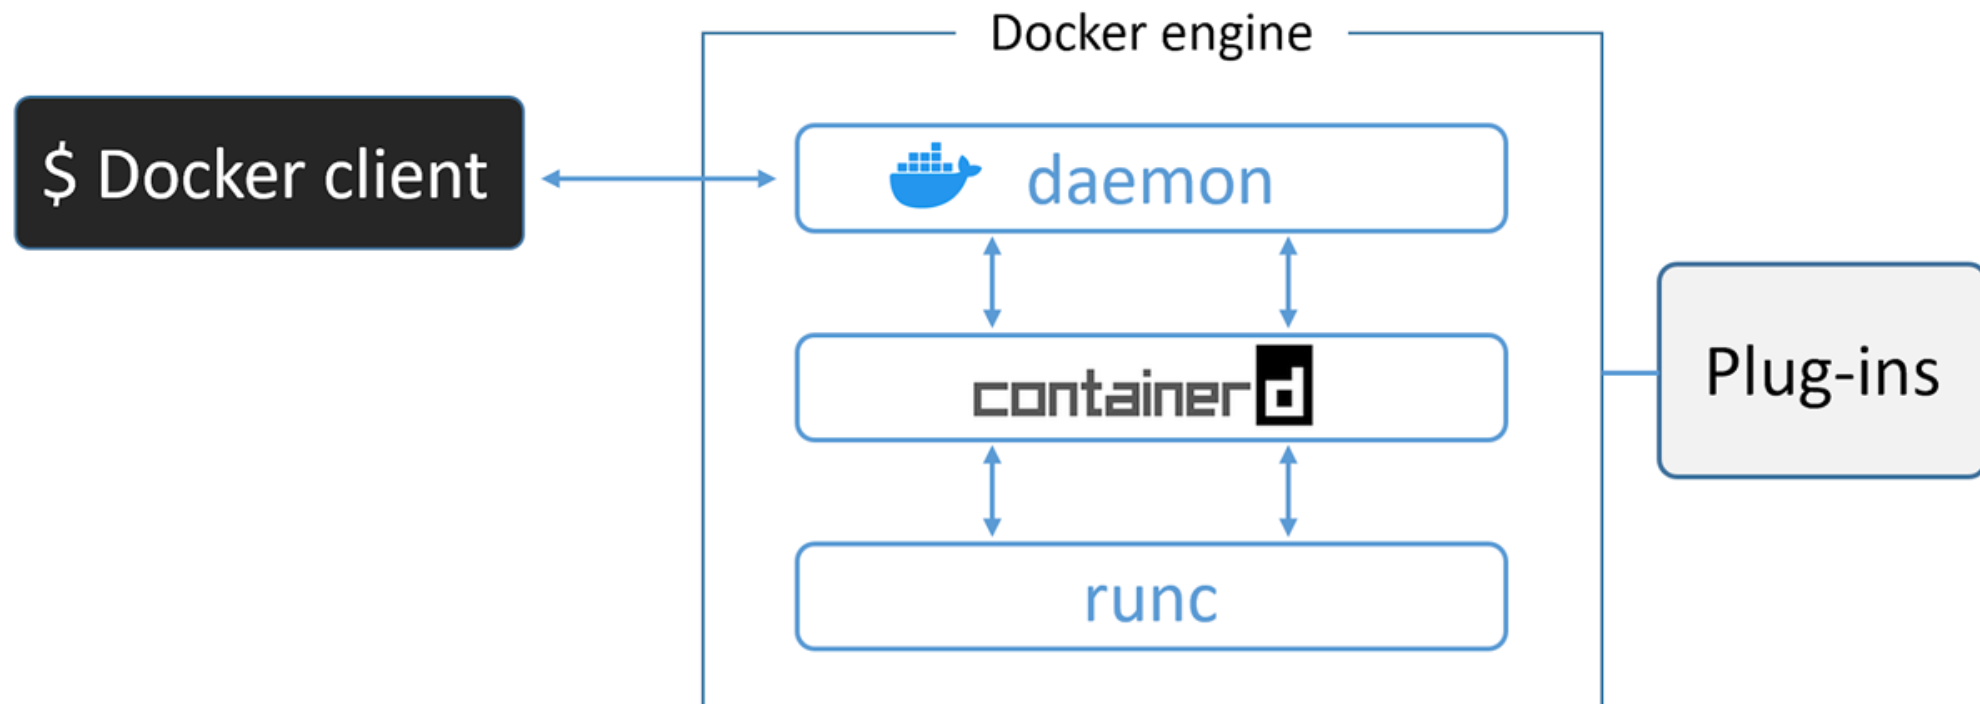
Developed in 2005 and introduced to containers circa 2014, Seccomp **limits** how processes could use **system calls**. Seccomp defines a security profile for processes that lists the system calls, parameters and file descriptors they are allowed to use.

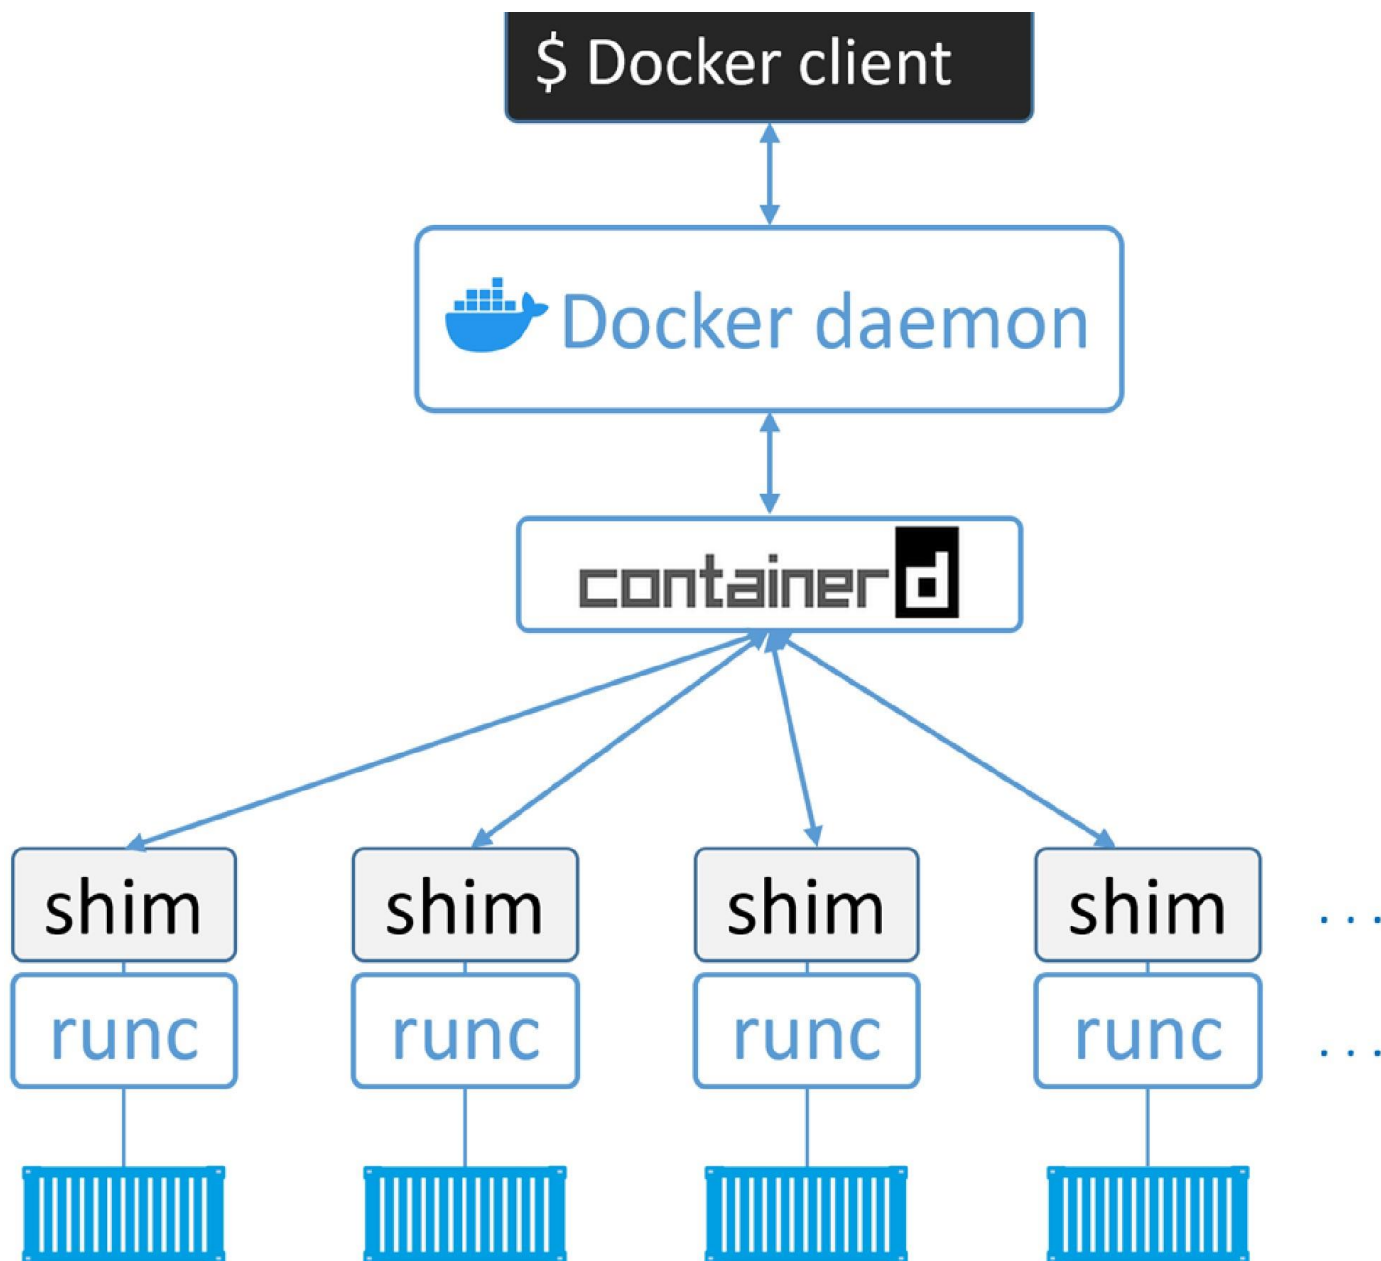
SELinux

Security-Enhanced Linux (SELinux) is a mandatory access control system for processes. Linux kernel uses SELinux to **protect processes** from **each other** and to protect the host system from its running processes. Processes run as a confined SELinux type that has limited access to host system resources



Docker architecture





Docker commands (CLI)



API and other features

Container supervisor
start | stop | pause ..

... Enables daemonless containers

... Container runtime
(interface to kernel primitives)

Running containers



\$ Docker client

Issue ``docker container run`` command
to **Docker API** exposed by Docker daemon

 Docker daemon

Receive instruction at API endpoint.
Instruct **containerd** (via gRPC API) to start new
container based on OCI bundle and ID provided

containerd

Receive instruction to create containers
Instruct **runc** to create container.

shim

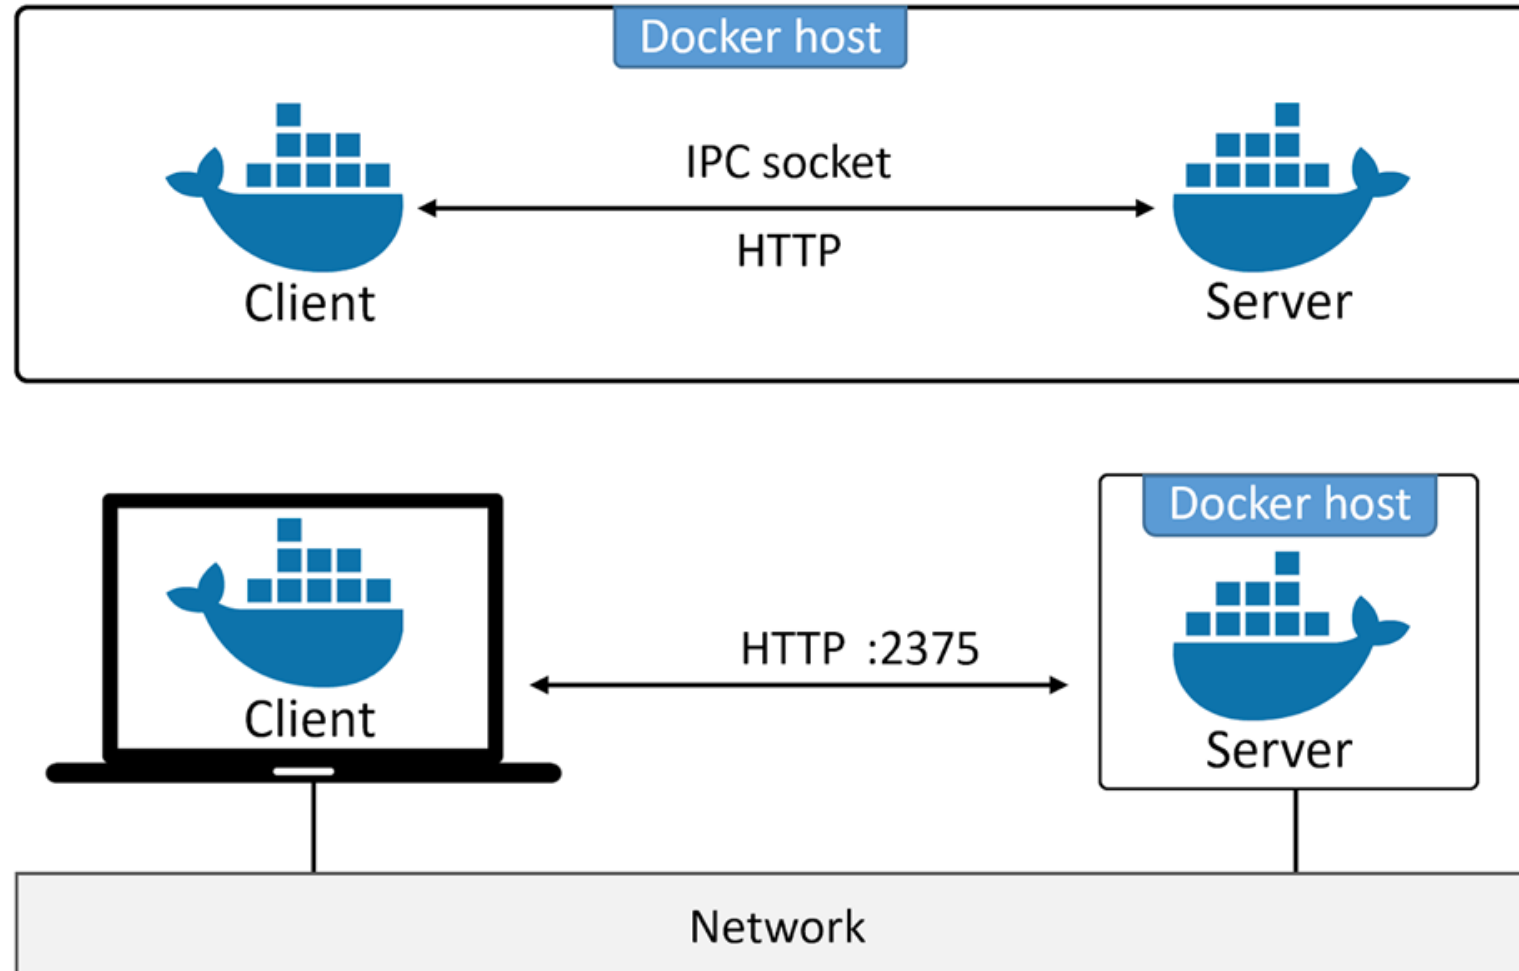
Build and start container
runc exit after container start
shim become container's parent process

runc






Securing client and daemon communication



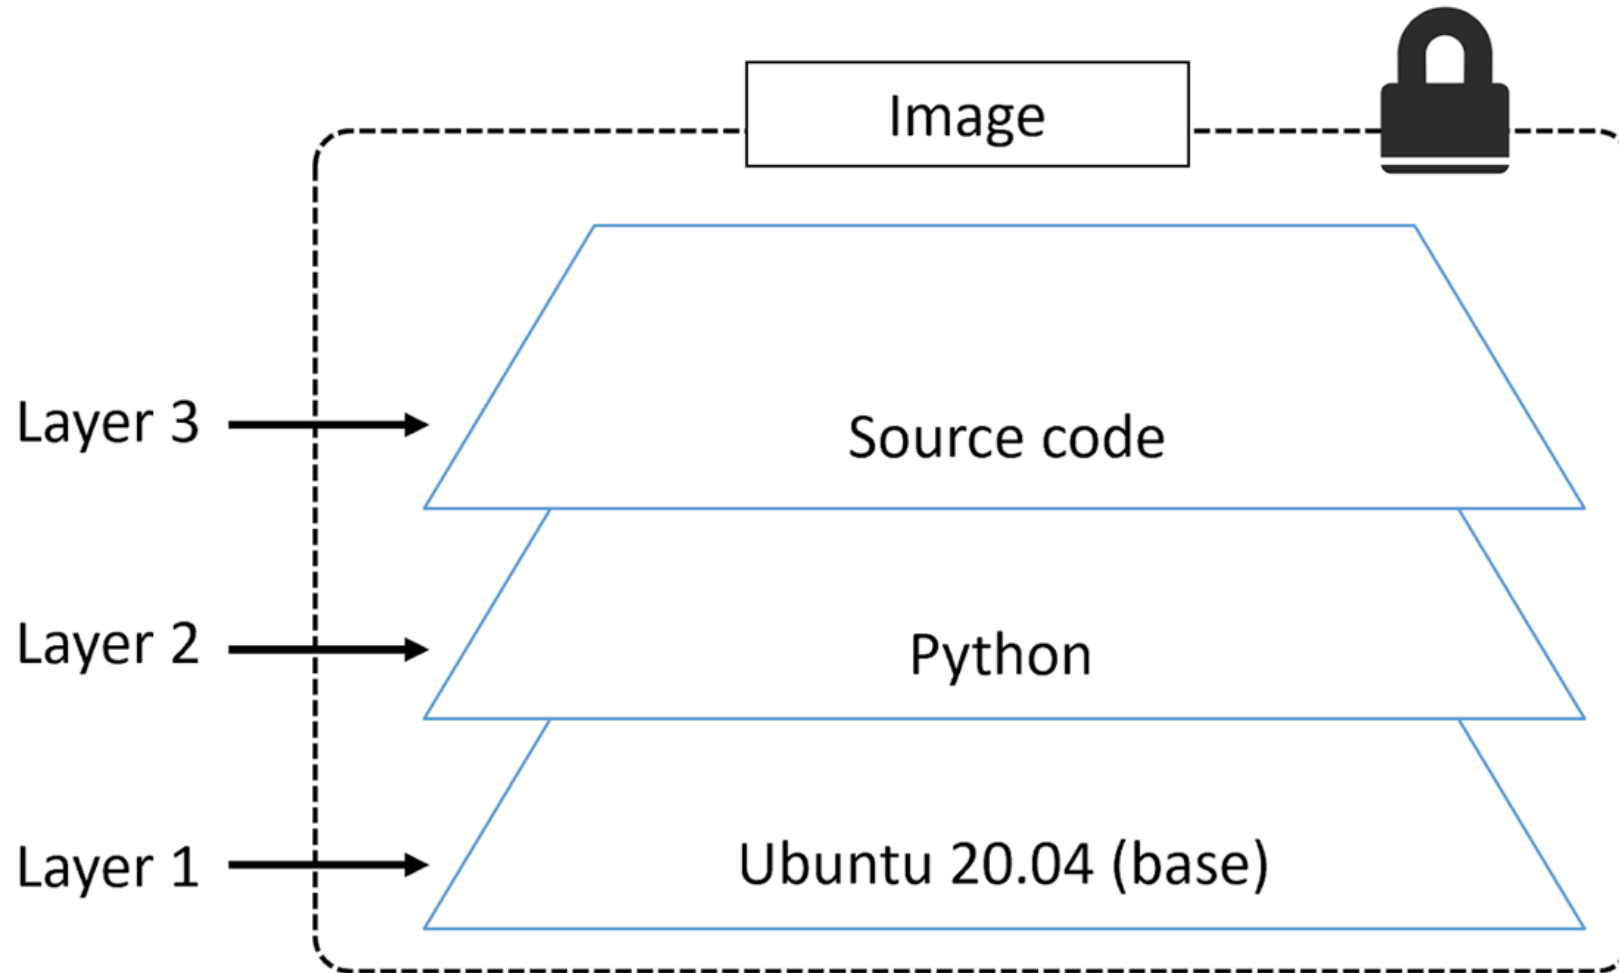


Docker Images

Read-only template that contains a set of instructions for creating a **container** that can run on the **Docker** platform.



Images and layers





Images and layers

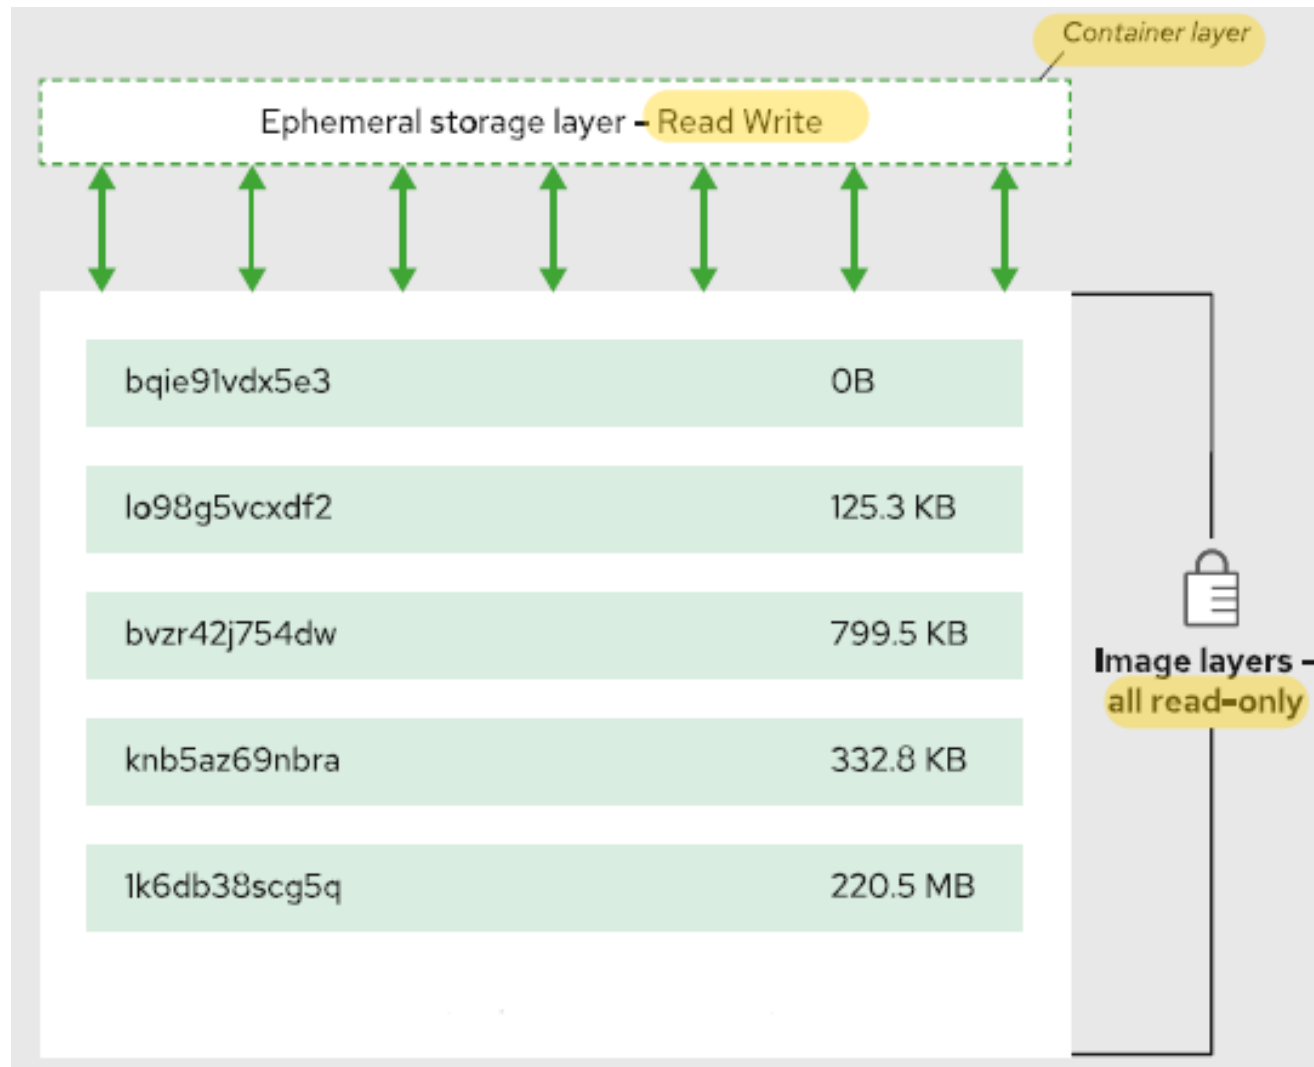
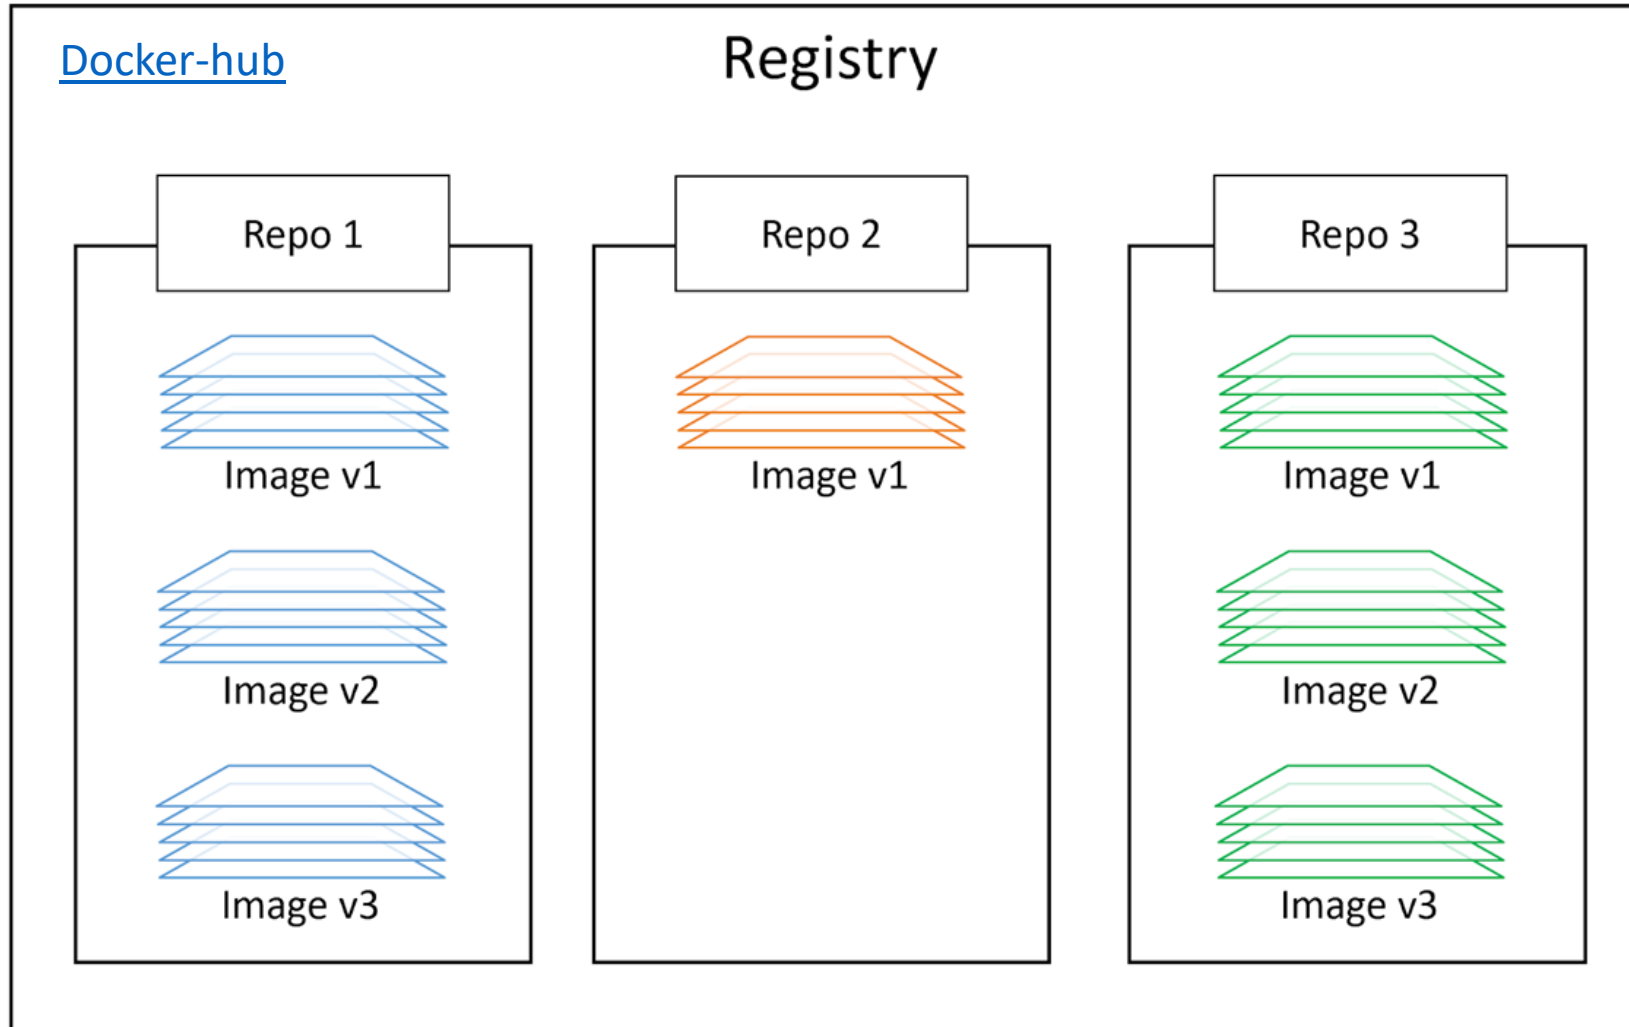



Image Registries





Docker Containers





Self-healing containers with restart policies

Always

- container will always restart if the **main process** is **killed** from inside the container but won't restart if you manually stopped it. Will restart if the Docker **daemon** restarts

unless-stopped

- container will always restart if the **main process** is **killed** from inside the container but won't restart if you manually stopped it. However will NOT restart if the Docker **daemon** restarts.

on-failure

- container will always restart if the **main process exits** with non-zero code (i.e. with error) but won't restart if you manually stopped it. However will restart if the Docker **daemon** restarts

Install Docker



[How To Install Docker](#)

[Docker-hub](#)



Images Commands

[docker images](#) shows all images.

[docker rmi](#) removes an image.

[docker search](#) image search

[docker history](#) shows history of image.

[docker tag](#) tags an image to a name (local or registry).



Images Commands cont'd

[docker commit](#) creates image from a container, pausing it temporarily if it is running.

Ex: `docker commit [OPTIONS] CONTAINER_ID [REPOSITORY[:TAG]]`

[docker save](#) saves **an image to a tar archive stream** to STDOUT with all parent layers, tags & versions

Ex : `sudo docker save busybox-1 > /home/save.tar`

[docker load](#) loads an **image from a tar archive** as STDIN, including images and tags (as of 0.7).

N.B: All the above commands will require the IMAGE_ID



Containers Commands

[docker start/stop/restart](#) starts/stops/restarts a container.

[docker run](#) creates and **starts a container in one operation.**

Ex: `docker run -p $HOSTPORT:$CONTAINERPORT --name CONTAINER some image`

Ex: **`docker run -it IMAGE bash`** (it will open the container)

[docker ps](#) shows running containers.

[docker ps -a](#) shows running and stopped containers.

[docker logs](#) gets logs from container.



Containers Commands cont'd

[docker inspect](#) looks at all the info on a container (including IP address).

[docker events](#) gets events from container.

[docker port](#) shows public facing port of container.

[docker top](#) shows running processes in container.

[docker stats](#) shows containers' resource usage statistics.

[docker diff](#) shows changed files in the container's FS.



Containers Commands cont'd

[docker export](#) – Exports a container's **filesystem as a tar archive**

Ex: `docker export <CONTAINER ID> > /home/export.tar`

[docker import](#) creates an **image from a tarball**.

[docker cp](#) copies local file to container and vice versa.

Ex: `docker cp [OPTIONS] CONTAINER_ID:SRC_PATH DEST_PATH`

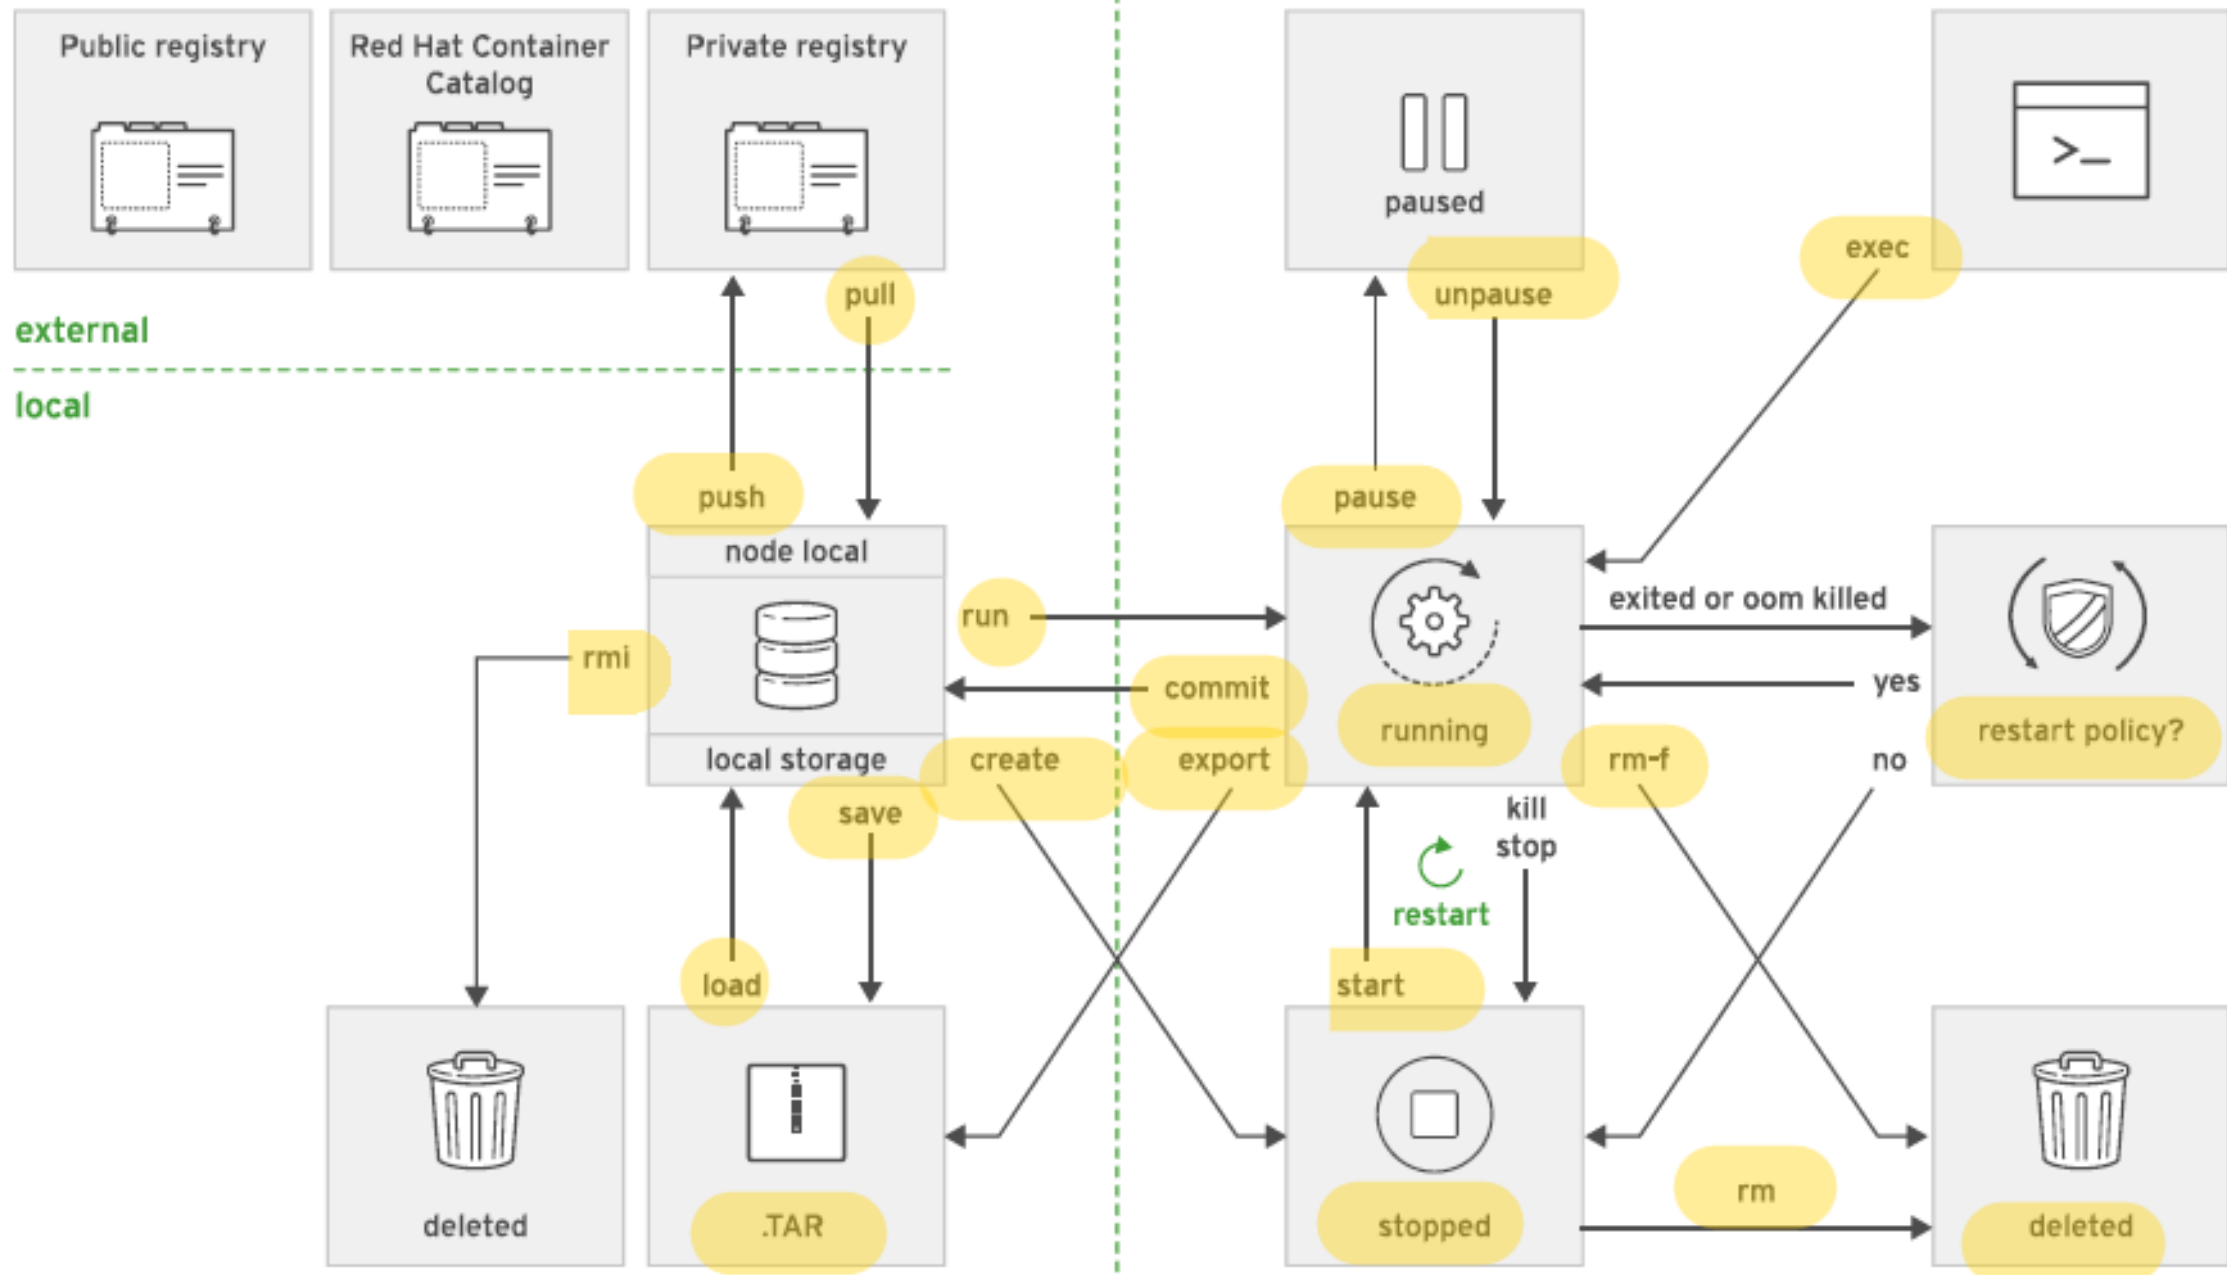
Ex: `docker cp [OPTIONS] SRC_PATH CONTAINER_ID:DEST_PATH`

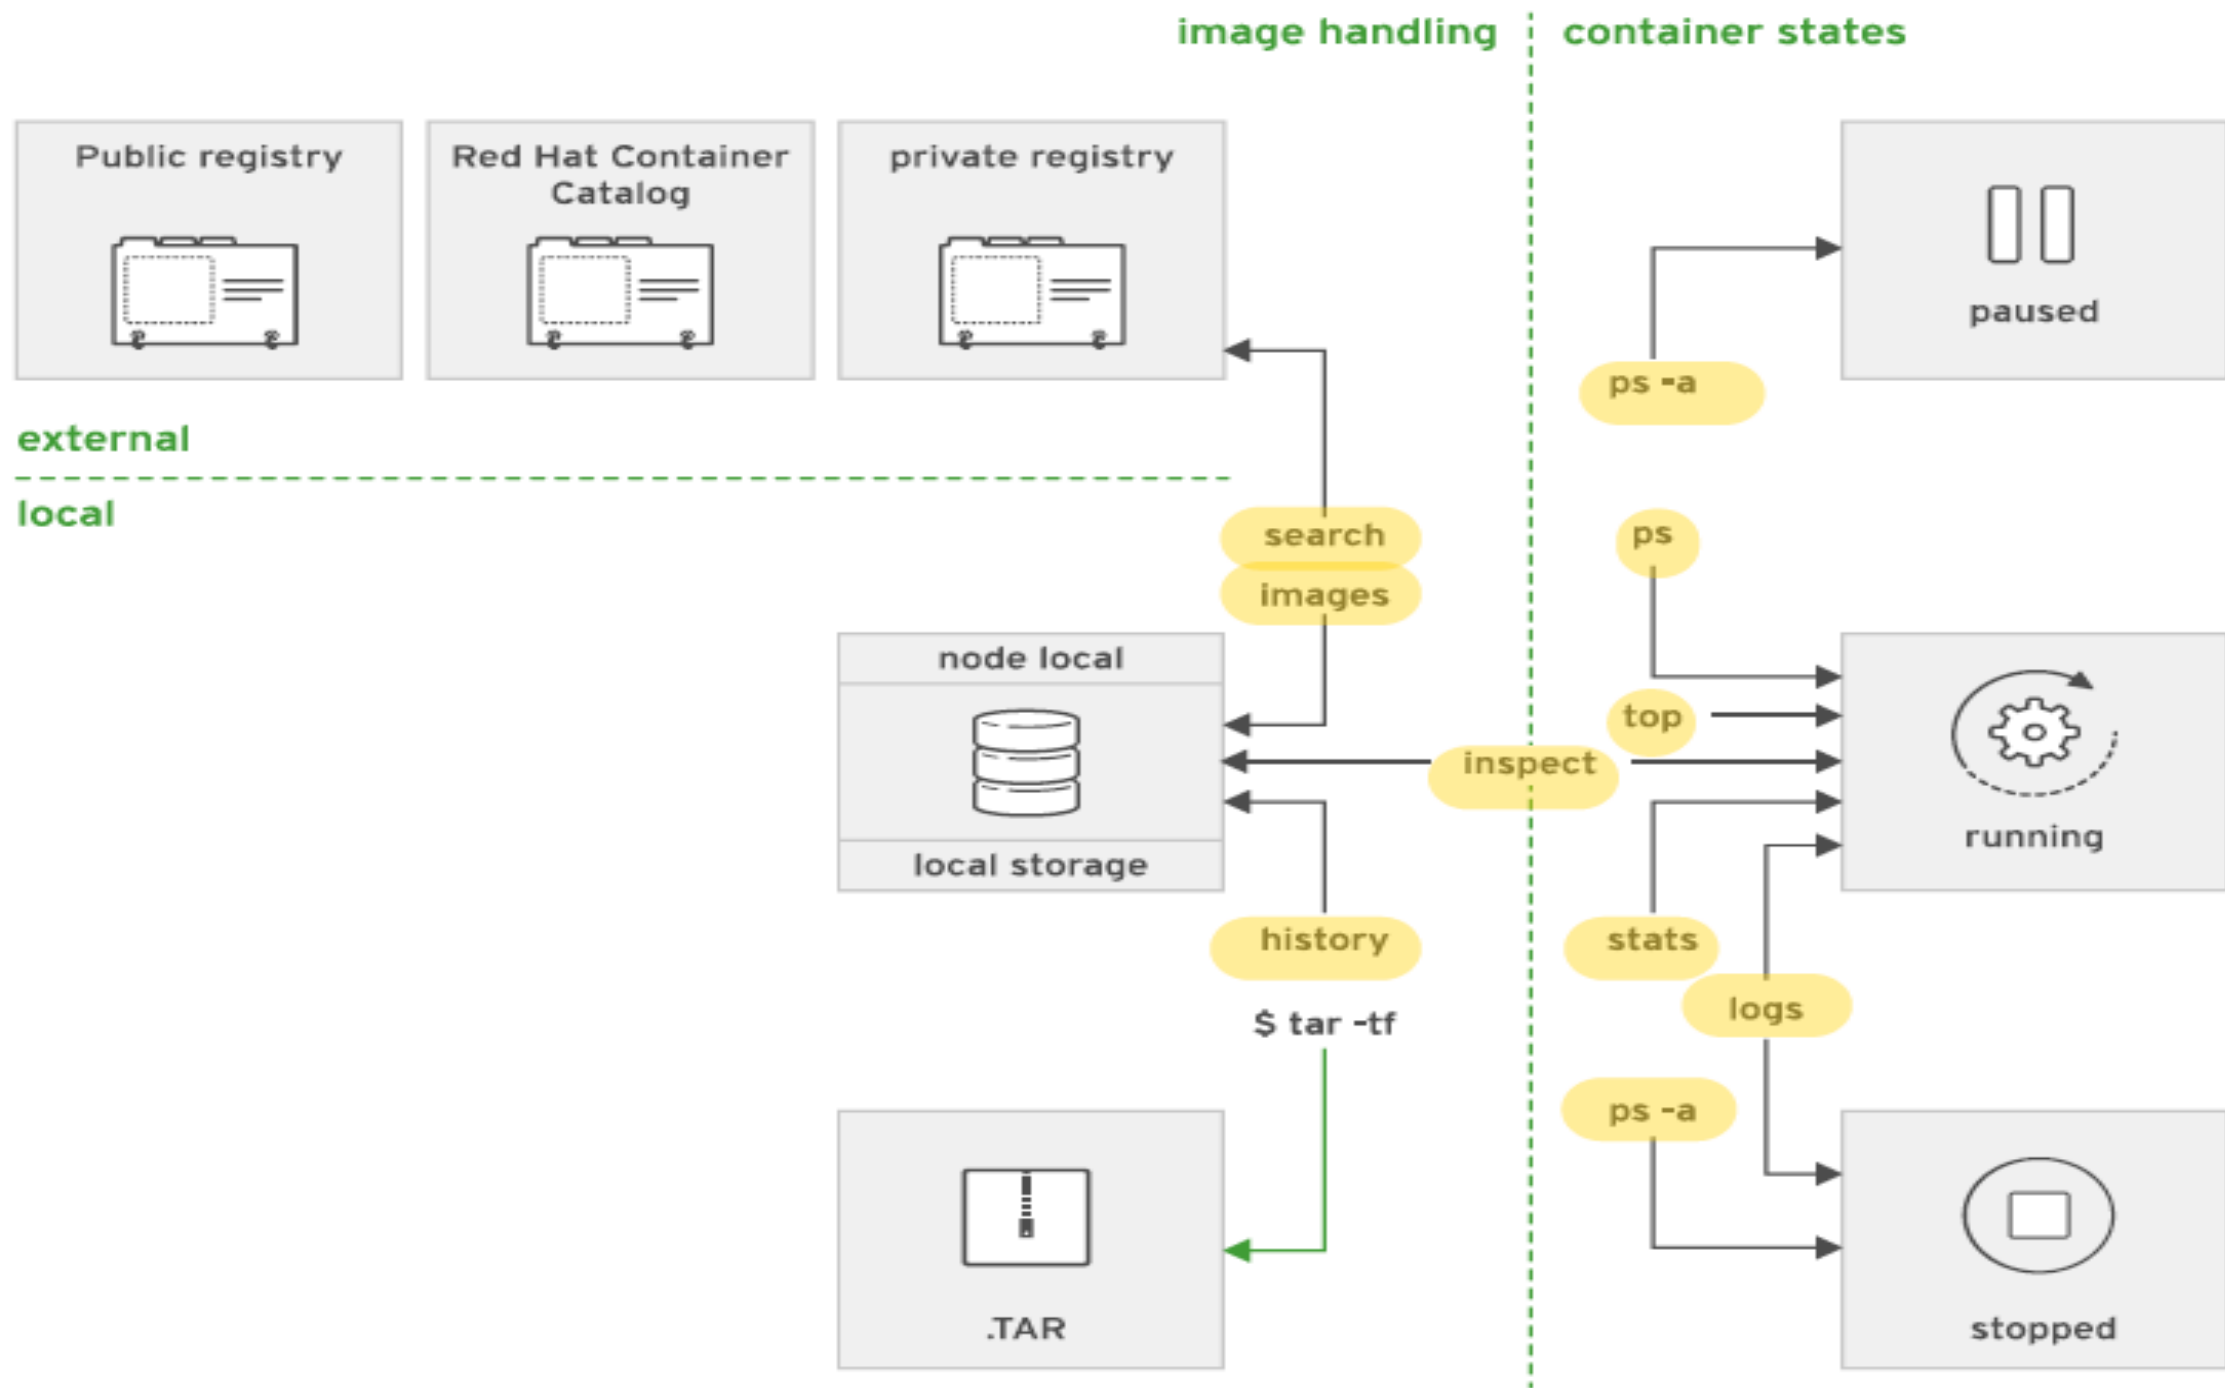
[docker exec](#) to execute a command inside a container.

Ex: `docker exec -it CONTAINER_ID bash`

N.B: All the above commands will require the CONTAINER_ID

image handling | container states







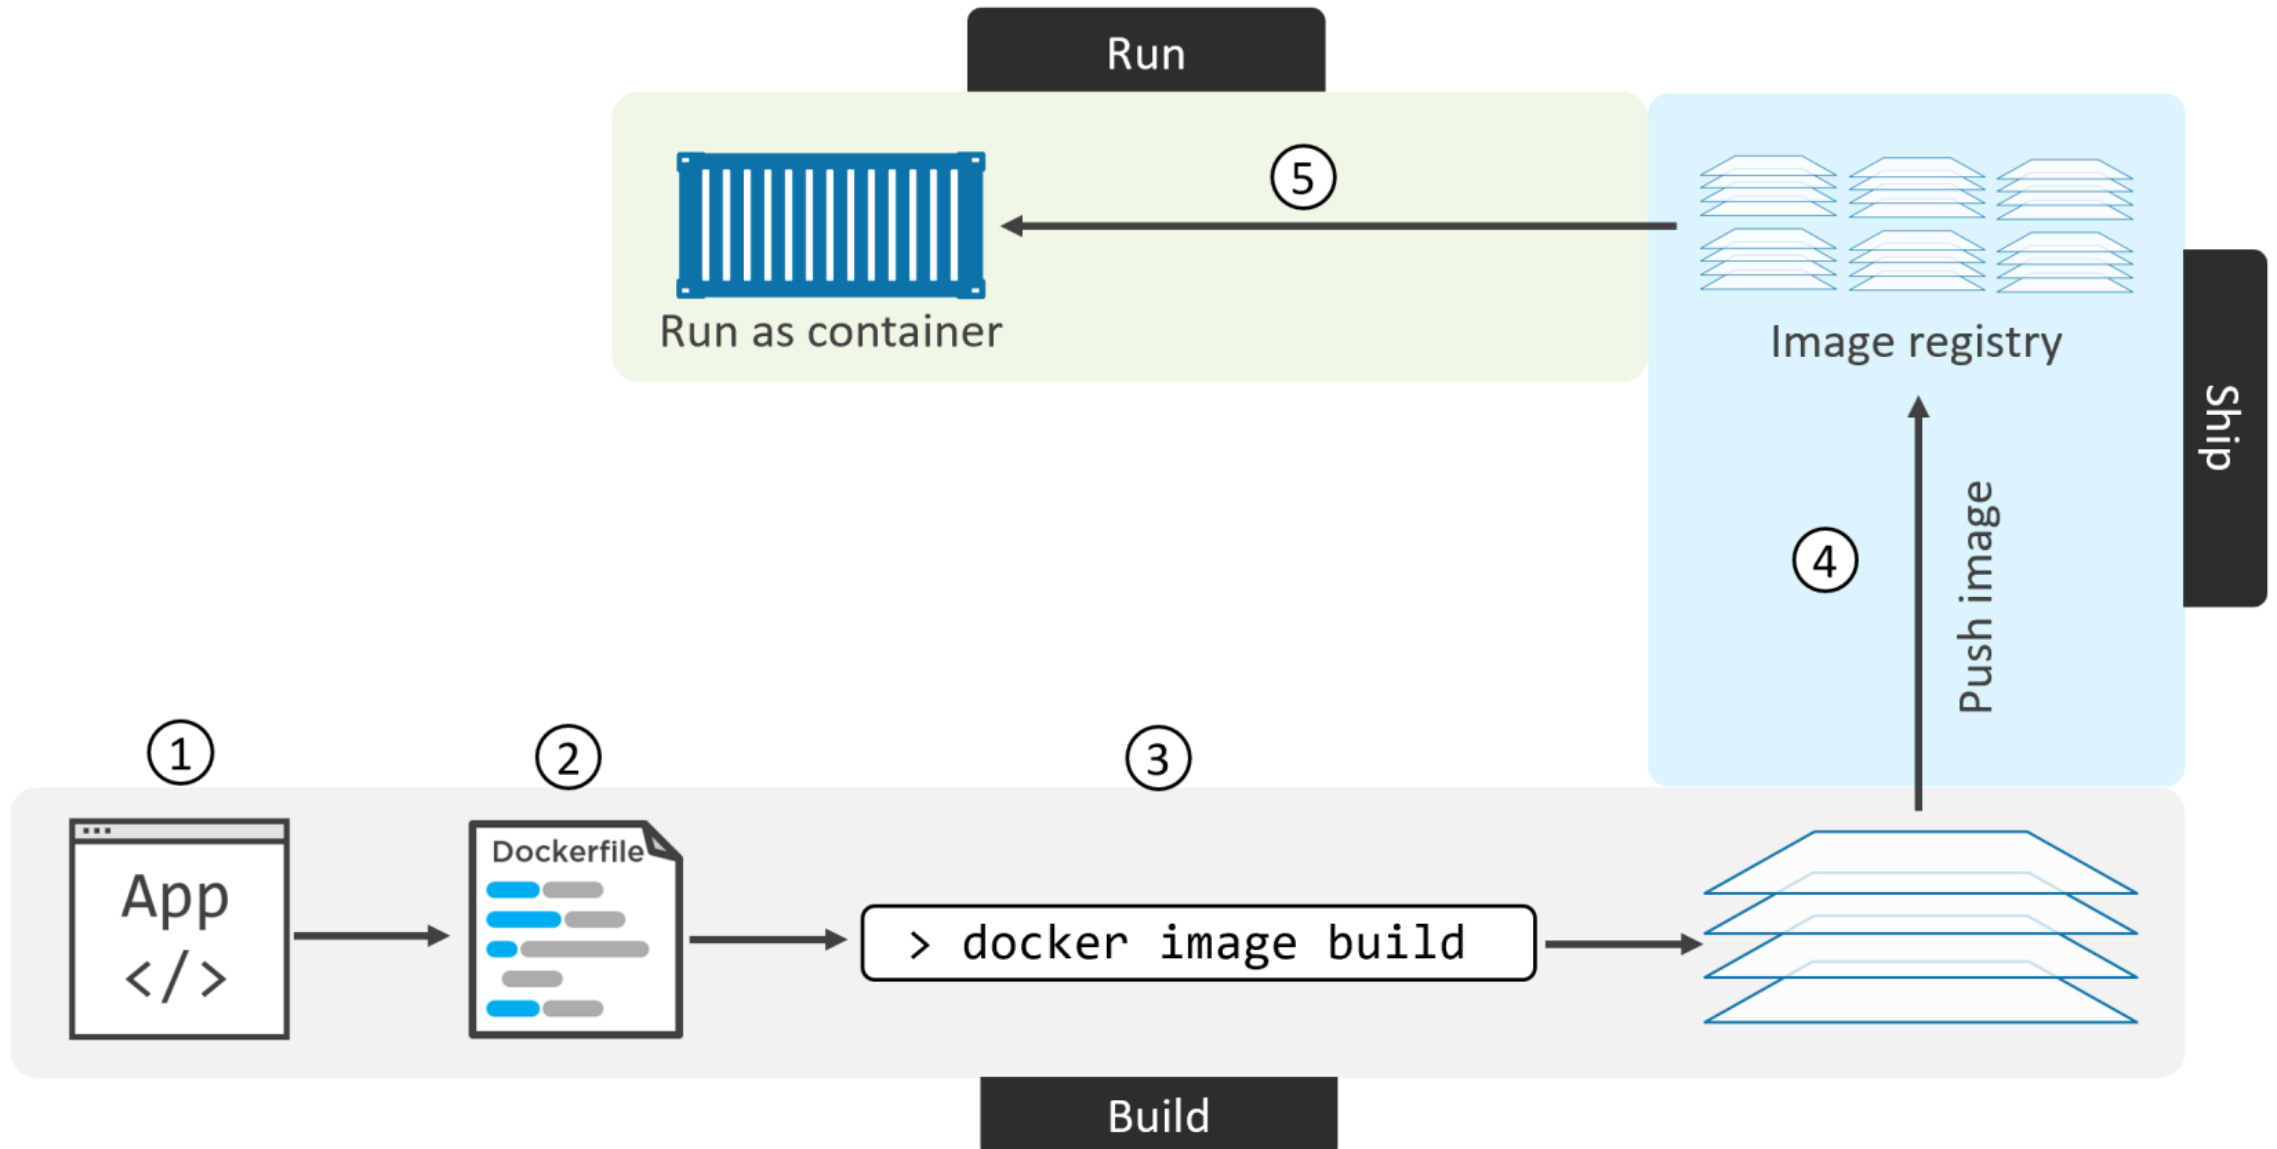
Docker file

Dockerfile is a script that contains instructions for building a customized docker image.

Each instruction in a Dockerfile creates a new layer in the image, and the final image is composed of all the layers stacked on top of each other

It includes instructions for installing dependencies, copying files, setting environment variables, and configuring the container

Docker file





Dockerfile basics

FROM Sets the Base **Image** for subsequent instructions, you should write this **command** at least in Dockerfile

RUN execute any commands in a new layer on top of the current image and commit the results. (Used mainly for installing packages)

CMD provide defaults for **executing command** inside a **container**.

EXPOSE informs Docker that the container listens on the specified network **ports** at runtime. NOTE: does not actually make ports accessible.



Dockerfile basics cont'd

ENV sets environment variable.

COPY copies **new files or directories to container** Note that this only copies as root, so you have to chown manually regardless of your USER / WORKDIR setting. See <https://github.com/moby/moby/issues/30110>

ADD lets you do that too, but it also supports 2 other sources. First, you can use **a URL instead** of a local file / directory. Secondly, you can **extract a tar** file from the source directly into the destination.



Dockerfile basics cont'd

ENTRYPOINT configures a container that will run as an executable = **CMD**

VOLUME creates a mount point for **externally mounted volumes** or other containers.

USER sets the user name for following RUN / CMD / ENTRYPOINT commands, By default it run commands as a **root**

WORKDIR sets the **working directory**, By default it run commands in **home directory**

ARG defines a build-time variable.



Dockerfile Sample

```
FROM node:14-alpine3.16
```

```
WORKDIR /app
```

```
COPY . .
```

```
RUN npm install
```

```
CMD [ "npm", "start" ]
```



Building Docker file, tagging and creating image

Traditionally, the Dockerfile is called "Dockerfile"

build Dockerfile > docker build .(current dictory)

build Dockerfile with image tag add "-t" > docker build -t "tag" .(current dictory)

If the Dockerfile is not named as "Dockerfile"

build Dockerfile > docker build -f /path/to/a/Dockerfile

build Dockerfile with image tag add "-t" > docker build -t "tag" -f /path/to/a/Dockerfile



Logging to your Docker repo from CLI

`docker login "URL"`

By Default, `docker login` > redirects login to docker hub

Enter your username and password



Pushing/Pulling image to docker repo

```
docker tag local-image:tagname remote-repo:tagname
```

```
docker push remote-repo:tagname
```

```
docker pull new-repo:tagname
```

```
docker tag    image-id hossamesaaa/firstproject:hossam
```

```
docker push  hossamesaaa/firstproject:hossam
```

```
docker pull  hossamesaaa/firstproject:hossam
```